

O'REILLY®



Live!

# Effective Modern C++


42 SPECIFIC WAYS TO IMPROVE YOUR USE OF C++11 AND C++14

Scott Meyers

# Consider void Futures for One-Shot Event Communication

Scenario:

- Concurrent tasks T1 and T2 begin.
- Among T1's work is some action A.
- At some point, T2 may continue only if A has occurred.



How have T1 tell T2 that A has taken place?

## Poll (Multi-Choice Single Answer)

How would you approach this problem?

- Use a condition variable: T1 notifies T2 when A has occurred.
- Use a flag: T2 polls a flag that T1 sets when A has occurred.
- Use a condition variable and a flag: Avoids some drawbacks of each approach above.
- Use a future: T2 blocks on a future that T1 sets.
- None of the above.

# Condition Variables

Basic idea:

- T2 waits on a condvar.
- T1 notifys condvar when A done.

```
std::condition_variable cv;  
std::mutex m;
```

```
...  
Do A  
cv.notify_one();  
...
```

Task T1

```
...  
std::unique_lock<std::mutex> lk(m);  
cv.wait(lk);  
...
```

Task T2

# Condition Variables

Three problems:

- Condvars require a mutex, but there's no state to protect.
  - ➔ "Feels wrong"
- Condvars may receive spurious notifications.
  - ➔ T2 must verify that A has occurred.
    - ◆ But that's our goal with the condvar!
- If T1 notifies before T2 waits, T2 blocks forever!
  - ➔ Must check to see if A has happened before waiting.
    - ◆ Again, our goal is to use the condvar for that.

```
...  
Do A  
cv.notify_one();  
...
```

Task T1

```
...  
std::unique_lock<std::mutex> lk(m);  
cv.wait(lk);  
...
```

Task T2

# std::atomic Flags

Basic idea:

- T1 sets a flag when A occurs.
- T2 polls flag and proceeds only when flag set.

```
std::atomic<bool> flag(false);
```

```
...  
Do A  
flag = true;  
...
```

Task T1

```
...  
while (!flag);  
...
```

Task T2

# `std::atomic` Flags

Problem:

- Polling keeps T2 running, even though it's conceptually blocked.
  - ➔ Uses CPU (hence power).
  - ➔ Prevents other threads from using its HW thread.
  - ➔ Incurs context switch overhead each time it's scheduled.

# Condition Variables + Flags

Basic idea:

- T2 waits on a condvar, checks a flag when notifyd.
- T1 sets a flag and notifies condvar when A occurs.

```
std::condition_variable cv;  
std::mutex m;  
bool flag;
```

```
...  
Do A  
{  
    std::lock_guard<std::mutex> g(m);  
    flag = true;  
}  
cv.notify_one();  
...
```

Task T1

```
...  
{  
    std::unique_lock<std::mutex> lk(m);  
    cv.wait(lk, []{ return flag; });  
    ...  
}  
...
```

Task T2



# Condition Variables + Flags

Problem:

- Stilted communication protocol.
  - ➔ T1 sets flag  $\Rightarrow$  A has occurred, but condvar notify still necessary.
  - ➔ T2 wakes  $\Rightarrow$  A has probably occurred, but flag check still necessary.

# void Futures

void is content-free. Its availability isn't.

- Unlike condvar:
  - ➔ No need for gratuitous mutex.
  - ➔ No need for “did event already/really occur?” mechanism.
- Unlike `std::atomic<bool>`, requires no polling.

# Example: Starting Thread Suspended

Motivation:

- Separate thread creation from running something on it.
- Available on some platforms (e.g. Windows)

No direct support in C++11.

- Easy to implement via `void future` (at least in concept...).
- Function to run on thread still required at thread creation.
  - ➔ “Function”  $\equiv$  “Callable object”

# Starting Thread Suspended


```
std::promise<void> p;
```

## Poll (Multi-Choice Single Answer)

How familiar are you with C++11's `std::promises`?

- I've written code using them.
- I haven't used them, but I've read about them.
- I've heard of them, but that's about it.
- I've never heard of them.

# Starting Thread Suspended



```
std::thread t2([&p]{ p.get_future().wait();    // start t2 and
                  funcToRun(); }             // suspend it
                  );                          // (conceptually)

...                                           // t2 is "suspended" waiting
                                           // for p to be set

p.set_value();                               // t2 may now continue

...

t2.join();
```

# `std::thread` Destructor Behavior

A `std::thread` object may be *joinable*:

- Represent an underlying (i.e., OS) thread of execution (TOE).

Unjoinable `std::threads` represent no underlying TOE.

- Default-constructed `std::threads`.
- `std::threads` that have been detached or moved from.
- `std::threads` whose TOE has been joined.
- Etc.

# std::thread Destructor Behavior

std::thread dtor calls std::terminate if it's joinable:

```
{  
    std::thread t(funcToRun);    // t is joinable  
    ...                          // assume t remains joinable  
}  
                                // std::terminate called
```



# `std::thread` Destructor Behavior

Implication: `std::threads` must be unjoinable on all paths from block:

- Be unjoinable:
  - ➔ Be joined or
  - ➔ Be detached or
  - ➔ Be otherwise made unjoinable (e.g., moved from)
- All paths:
  - ➔ `continue`, `break`, `goto`, `return`
  - ➔ Flow off end
  - ➔ Exception

# Unjoinable on *All Paths*

All paths  $\Rightarrow$  RAII classes.

- None for `std::thread` in standard library!

- ➔ From the Standard:

- Either implicitly detaching or joining a `joinable()` thread in its destructor could result in difficult to debug correctness (for detach) or performance (for join) bugs encountered only when an exception is raised.

- Easy to write your own.

# RAII Class for `std::thread`

```
class ThreadRAII {
public:
    enum class DtorAction { join, detach };

    ThreadRAII(std::thread&& t, DtorAction a)    // in dtor, take
    : action(a), t(std::move(t)) {}           // action a on t

    ~ThreadRAII()
    {
        if (t.joinable()) {
            if (action == DtorAction::join) t.join();
            else t.detach();
        }
    }

    std::thread& get() { return t; }

private:
    DtorAction action;
    std::thread t;
};
```

# RAII Class for `std::thread`

```
ThreadRAII t1(std::thread(doThisWork),           // join on  
              ThreadRAII::DtorAction::join);      // destruction  
  
ThreadRAII t2(std::thread(doThatWork),           // detach on  
              ThreadRAII::DtorAction::detach);    // destruction
```

# RAII Class for `std::thread`

Dtor prevents generation of move ops, but they make sense, so:

```
class ThreadRAII {
public:
    enum class DtorAction { join, detach };           // as before

    ThreadRAII(std::thread&& t, DtorAction a)           // as before
    : action(a), t(std::move(t)) {}

    ~ThreadRAII()
    {
        ... // as before
    }

    ThreadRAII(ThreadRAII&&) = default;                // support
    ThreadRAII& operator=(ThreadRAII&&) = default;     // moving

    std::thread& get() { return t; }                  // as before

private:                                              // as before
    DtorAction action;
    std::thread t;
};
```

# Starting Thread Suspended

Use of ThreadRAII ensures `join` is always performed:

```
std::promise<void> p;                                // as before
std::thread t2([&p]{ p.get_future().wait(); // as before
                  funcToRun(); }
               );
ThreadRAII tr(std::move(t2), ThreadRAII::DtorAction::join);
...                                                    // as before

p.set_value();                                         // as before
...
t2.join();                                             // no longer needed
```

# Poll (Multi-Choice Multi-Answer)

Which of these consequences arise from using ThreadRAII?

- The ThreadRAII object will occupy stack space.
- ThreadRAII's ctor/dtor will reduce the speed of the code.
- The code is exception safe.

# Starting Thread Suspended

Use of ThreadRAII ensures `join` is always performed:

```
std::promise<void> p;           // as before
std::thread t2([&p]{ p.get_future().wait(); // as before
                  funcToRun(); }
               );
```

```
ThreadRAII tr(std::move(t2), ThreadRAII::DtorAction::join);
```

```
...
```

```
p.set_value();
```

```
...
```


```
t2.join();
```

```
// handled by ThreadRAII dtor
```



# When A Doesn't Occur

What if T1 needs to tell T2 that A didn't occur (e.g., due to an exception)?



- **Condvar approach:** No `notify`  $\Rightarrow$  T2 blocks forever (modulo spurious wakes).
- **Flag approach:** Flag never set  $\Rightarrow$  T2 polls forever.
- **Condvar + Flag approach:** No `notify` + flag never set  $\Rightarrow$  T2 blocks forever (maybe with occasional spurious wakes).

Must use out-of-band mechanism to communicate “A didn't occur”.

# When A Doesn't Occur

Observations:

- Futures can hold exceptions.
- `std::promise` ctor puts an exception into an unset promise.

Ergo:

- Use local `std::promise` in T1.
  - ➔ It's set  $\Rightarrow$  T2 "receives" void.
  - ➔ It's destroyed w/o being set  $\Rightarrow$  T2 receives an exception.
- In T2, use `get` instead of `wait`.
  - ➔ Allows T2 to determine whether A occurred.
    - ◆ Non-exception: it did.
    - ◆ Exception: it didn't.

# Starting Thread Suspended

```
{ // make p local
    std::promise<void> p;
    std::thread t2([&p]
    {
        try {
            p.get_future().get();
            funcToRun(); // A occurred
        }
        catch(...) {
            ... // A didn't occur
        }
    });
    ThreadRAII tr(std::move(t2), ThreadRAII::DtorAction::join);
    ...
    p.set_value();
    ...
} // exception written to p
// if it hasn't been set
```

## Poll (Multi-Choice Single Answer)

Does this work, i.e., ensure that T2 gets an exception if A doesn't occur in T1?

- Yes.
- No.
- Leave me alone—it's nearly 5AM in Sydney, and it's already tomorrow!

# Starting Thread Suspended

```
{
    std::promise<void> p;                // created first,
                                        // destroyed last

    std::thread t2([&p]
    {
        try {
            p.get_future().get();
            funcToRun();
        }
        catch(...) { ... }
    });

    ThreadRAII tr(std::move(t2),         // created after
                  ThreadRAII::DtorAction::join); // p, destroyed
                                                // before p
    ...
    p.set_value();
    ...
}                                         // if p hasn't been set, tr's
                                         // dtor hangs on a join
```

# A Clean, Simple, Straightforward, Non-Error-Prone Solution

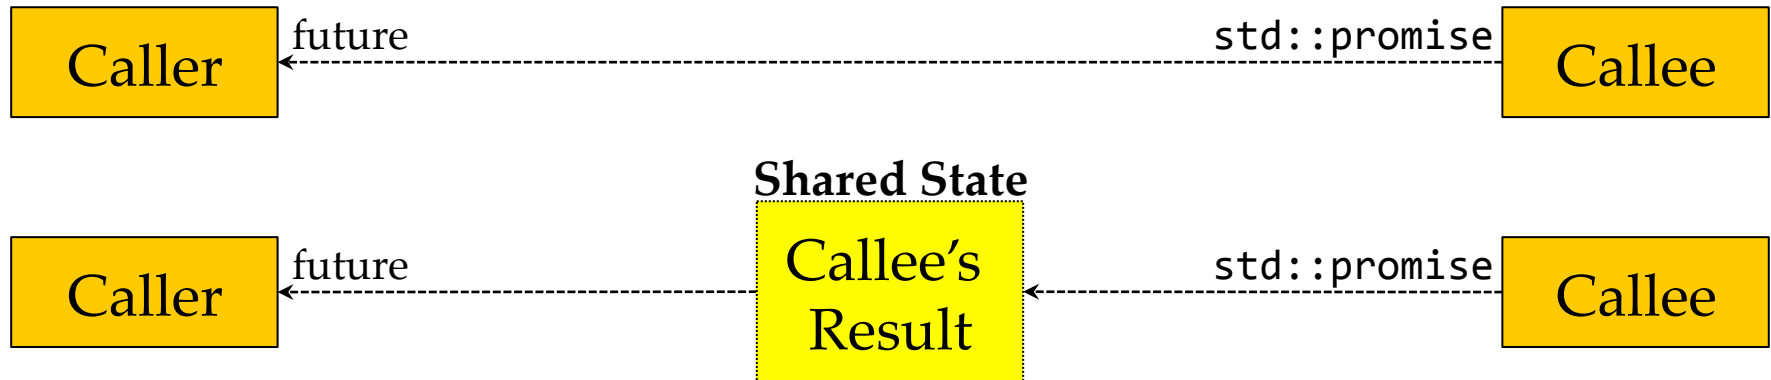
I don't know of one.

- I'm sorry, too.



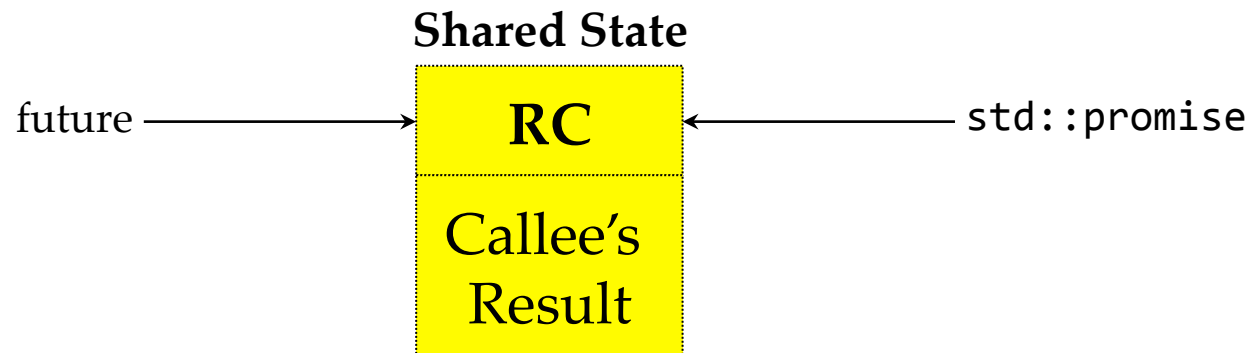
Image: "Depressed puma?," (c) Tambako The Jaguar @ Flickr. License: Creative Commons Attribution-NoDerivs 2.0 Generic.

# Costs for void Futures



Shared state:

- **Dynamically allocated**  $\Rightarrow$  use of heap.
- **Reference-counted**  $\Rightarrow$  atomic increments/decrements.



# *One-Shot* Event Communication

Shared state may be written at most *once*.

- Not suitable for recurring communication.

Alternatives avoid this restriction:

- Condition variables.
  - ➔ May be notify-ed repeatedly.
- `std::atomics`.
  - ➔ May be assigned repeatedly.




# Guideline

Consider `void` futures for one-shot event communication.

# Further Information

Source for this talk:


- *Effective Modern C++*, Scott Meyers, O'Reilly, 2015.
  - ➔ Item 17: Special member function generation.
  - ➔ Item 37: ThreadRAII.
  - ➔ Item 38: shared state.
  - ➔ Item 39: void futures.
- Email from Tomasz Kamiński.




# Further Information

The C++11 Concurrency API:

- *The C++ Standard Library, Second Edition*, Nicolai M. Josuttis, Addison-Wesley, 2012.



- *C++ Concurrency in Action*, Anthony Williams, Manning, 2012.



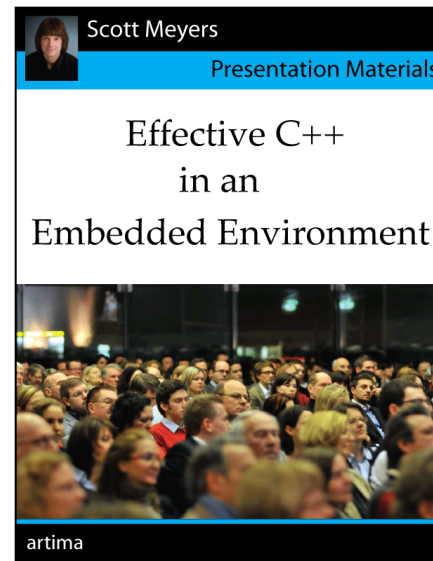
- “Why would I want to start a thread ‘suspended’?,” *Stack Overflow*, asked 1 July 2010.
- “ThreadRAII + Thread Suspension = Trouble?,” Scott Meyers, *The View from Aristeia*, 24 December 2013.
- “std::futures from std::async aren't special!,” Scott Meyers, *The View from Aristeia*, 20 March 2013.

# Licensing Information

Scott Meyers licenses materials for this and other training courses for commercial or personal use. Details:

- **Commercial use:** <http://aristeia.com/Licensing/licensing.html>
- **Personal use:** <http://aristeia.com/Licensing/personalUse.html>

Courses currently available for personal use include:



# About Scott Meyers



Scott offers training and consulting services on the design and implementation of C++ software systems. His web site,

<http://aristeia.com/>

provides information on:

- Training and consulting services
- Books, articles, other publications
- Upcoming presentations
- Professional activities blog